

Dr Renata Borovica-Gajic
David Eccles



INFO90002 Database Systems & Information Modelling

Lecture 09
SQL



Find all pairs of sailors in which the older sailor has a lower rating

Boats

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

**Reserves
(R1)**

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

**Sailors 1
(S1)**

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

**Sailors 2
(S2)**

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



- Find all pairs of sailors in which the older sailor has a lower rating

$$\rho (S1(1 \rightarrow sid1, 2 \rightarrow sname1, 3 \rightarrow rating1, 4 \rightarrow age1), Sailors)$$
$$\rho (S2(1 \rightarrow sid2, 2 \rightarrow sname2, 3 \rightarrow rating2, 4 \rightarrow age2), Sailors)$$
$$\pi_{sname1, sname2}(S1 \bowtie_{age1 > age2 \wedge rating1 < rating2} S2)$$



SQL – Part 1

Introduction to DDL, DML & Referential Integrity



- SQL – or SEQUEL is a language used in relational databases
- DBMS support CRUD
 - Create, Read, Update, Delete commands
- SQL supports CRUD
 - Create, Select, Update, Delete commands
- Other info
 - You can see the 2011 standard of SQL at
 - http://www.jtc1sc32.org/doc/N2151-2200/32N2153T-text_for_ballot-FDIS_9075-1.pdf
 - Wikipedia has several sections on SQL (good for generic syntax)
 - http://en.wikipedia.org/wiki/Category:SQL_keywords



- Provides the following capabilities:
 - Data Definition Language (DDL)
 - To define and set up the database
 - CREATE, ALTER, DROP
 - Data Manipulation Language (DML)
 - To maintain and use the database
 - SELECT, INSERT, DELETE, UPDATE
 - Data Control Language (DCL)
 - To control access to the database
 - GRANT, REVOKE
 - Other Commands
 - Administer the database
 - Transaction Control
 - START TRANSACTION
 - BEGIN, END



- In **Implementation** of the database
 - Take the tables we design in physical design
 - Implement these tables in the database using create commands
- In **Use** of the database
 - Use Select commands to read the data from the tables, link the tables together etc
 - Use alter, drop commands to update the database
 - Use insert, update, delete commands to change data in the database



1.

```
CREATE TABLE BankHQ (
    BankHQID INT AUTO_INCREMENT,
    HQAddress VARCHAR(120) NOT NULL,
    OtherHQDetails VARCHAR(100),
    PRIMARY KEY(BankHQID)
);
```

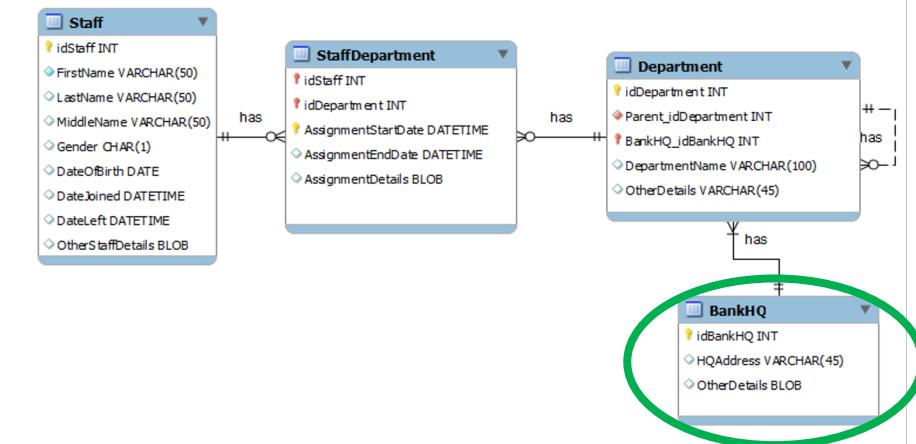
2.

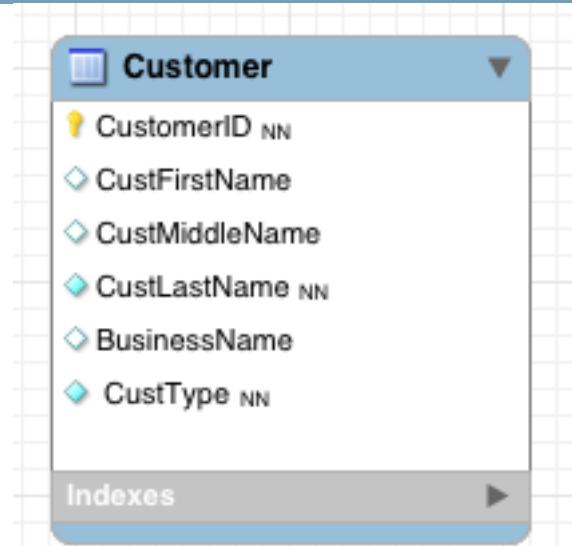
```
INSERT INTO BankHQ VALUES
(1, "23 Charles St Peterson North 2022", "Main Branch");
INSERT INTO BankHQ VALUES
(2, "213 Jones Rd Parkville North 2122", "Sub Branch");
```

3.

```
SELECT *
FROM BankHQ;
```

BankHQID	HQAddress	OtherHQDetails
1	23 Charles St Peterson North 2022	Main Branch
2	213 Jones Rd Parkville North 2122	Sub Branch

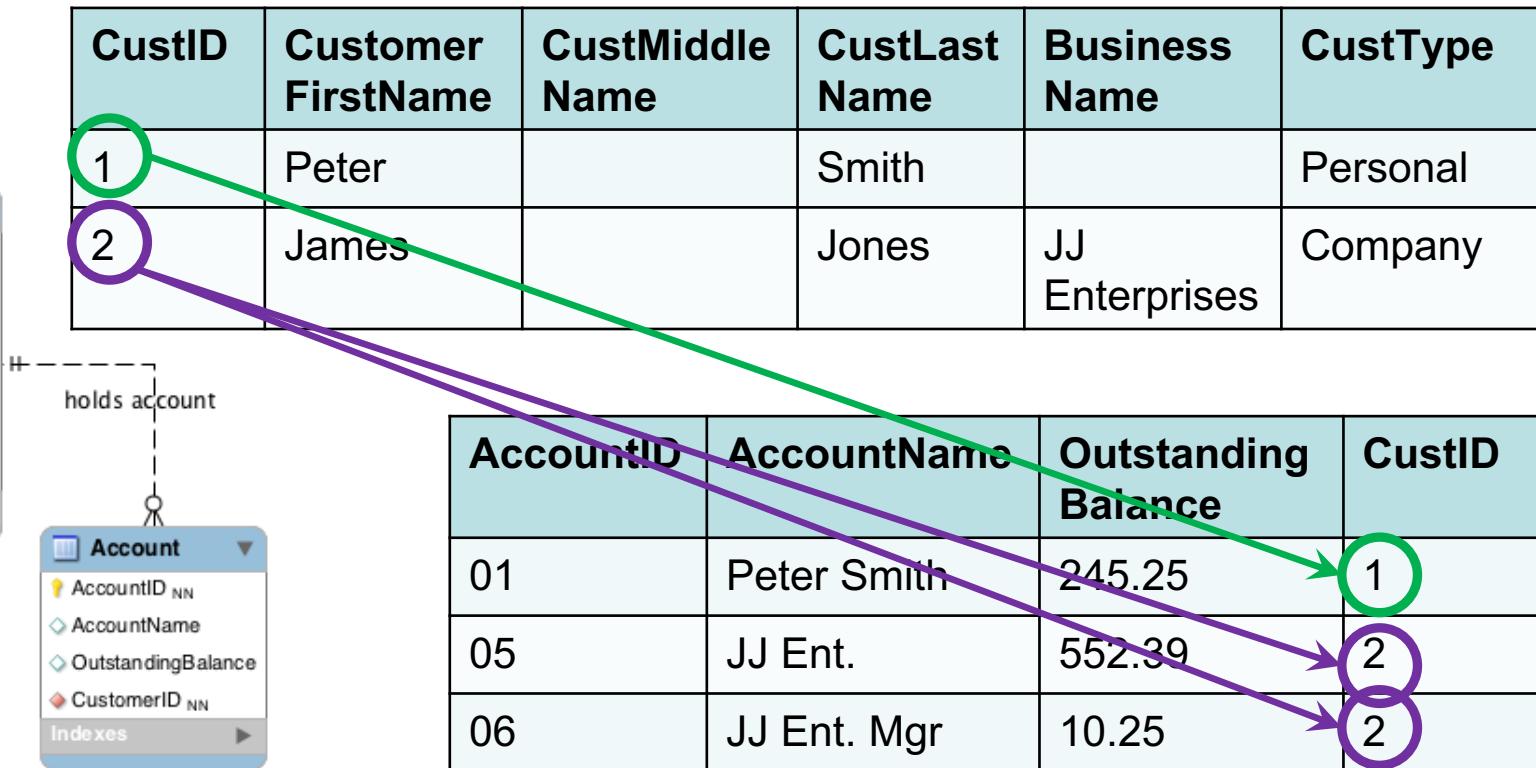




```
CREATE TABLE Customer (
    CustomerID          INT              AUTO_INCREMENT,
    CustFirstName       VARCHAR(12),
    CustMiddleName      VARCHAR(14),
    CustLatName         VARCHAR(20)        NOT NULL,
    Business Name       VARCHAR(100),
    CustType            ENUM("Personal", "Company") NOT NULL,
    PRIMARY KEY(CustomerID)
)
```



- We looked at Customer
 - A customer can have a number of Accounts
 - The tables get linked through a foreign key





```
CREATE TABLE Account (
    AccountID          INT             AUTO_INCREMENT,
    AccountName        VARCHAR(12),
    OutstandingBalance DECIMAL(10,2)      NOT NULL,
    CustomerID         VARCHAR(20)       NOT NULL,
    PRIMARY KEY(AccountID),
    FOREIGN KEY (CustomerID) REFERENCES Customer(CustomerID)
        ON DELETE RESTRICT
        ON UPDATE CASCADE
)
```



```
INSERT INTO Customer (CustomerID, CustFirstName, CustLastName, CustType)
    VALUES (DEFAULT, "Peter", "Smith", "Personal");
```

Specifies which columns will be entered

```
INSERT INTO Customer
    VALUES (DEFAULT, "James", NULL, "Jones", "JJ Enterprises", "Company");
```

```
INSERT INTO Customer
    VALUES (DEFAULT, "", NULL, "Smythe", "", "Company");
```

No column specification means ALL columns need to be entered

Customer

CustID	CustomerFirstName	CustMiddleName	CustLastName	BusinessName	CustType
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3		NULL	Smythe		Company

Null Island: The Busiest Place That Doesn't Exist:
<https://www.youtube.com/watch?v=bjvlpl-1w84>
by the channel MinuteEarth



- A cut down version of the SELECT statement – MySQL
- **SELECT [ALL | DISTINCT] *select_expr* [, *select_expr* ...]**
 - List the columns (and expressions) that are returned from the query
- **[FROM *table_references*]**
 - Indicate the table(s) or view(s) from where the data is obtained
- **[WHERE *where_condition*]**
 - Indicate the conditions on whether a particular row will be in the result
- **[GROUP BY {*col_name* | *expr*} [ASC | DESC], ...]**
 - Indicate categorisation of results
- **[HAVING *where_condition*]**
 - Indicate the conditions under which a particular category (group) is included in the result
- **[ORDER BY {*col_name* | *expr* | *position*} [ASC | DESC], ...]**
 - Sort the result based on the criteria
- **[LIMIT {[*offset*,] *row_count* | *row_count* OFFSET *offset*}]**
 - Limit which rows are returned by their return order (ie 5 rows, 5 rows from row 2)

Order is important! E.g. LIMIT cannot go before GROUP BY or HAVING



SQL

```
SELECT *
FROM Customer
```

RESULT

	CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType
▶	1	Peter	NULL	Smith	NULL	Personal
	2	James	NULL	Jones	JJ Enterprises	Company
	3	Akin	NULL	Smithies	Bay Wart	Company
	4	Julie	Anne	Smythe	Konks	Company
	5	Jen	NULL	Smart	BRU	Company
	6	Lim	NULL	Lam	NULL	Personal
	7	Kim	NULL	Unila	Saps	Company
	8	James	Jay	Jones	JJ's	Company
	9	Keith	NULL	Samson	NULL	Personal



The diagram shows a table structure for 'Customer'. The columns listed are CustomerID (PK), CustFirstName, CustMiddleName, CustLastName (PK), BusinessName, and CustType.

Customer	
CustomerID	NN
CustFirstName	
CustMiddleName	
CustLastName	NN
BusinessName	
CustType	NN

Indexes

In Relational Algebra:

$$\pi_{CustLastName}(Customer)$$

In SQL:

```
SELECT CustLastName  
FROM Customer;
```

Result

CustLastName
Smith
Jones
Smithies
Smythe
Smart
Lam
Unila
Jones
Samson

NOTE: MySQL doesn't discard duplicates.
To remove them use DISTINCT in front of
the projection list.



In Relational Algebra:

$$\sigma_{cond1 \wedge cond2 \vee cond3}^{(Rel)}$$

In SQL:

WHERE cond1 AND cond2
OR cond3

In Relational Algebra:

$$\pi_{CustLastName}(\sigma_{CustLastName = "Smith"}(Customer))$$

In SQL:

SELECT CustLastName
FROM Customer
WHERE CustLastName = "Smith";

Result

CustLastName
Smith



- In addition to arithmetic expressions, string conditions are specified with the LIKE clause

LIKE "REG_EXP"

% Represents zero, one, or multiple characters
_ Represents a single character

Examples:

WHERE CustomerName LIKE 'a%'	Finds any values that start with "a"
WHERE CustomerName LIKE '%a'	Finds any values that end with "a"
WHERE CustomerName LIKE '%or%'	Finds any values that have "or" in any position
WHERE CustomerName LIKE '_r%'	Finds any values that have "r" in the second position
WHERE CustomerName LIKE 'a_%_%'	Finds any values that start with "a" and are at least 3 characters in length
WHERE ContactName LIKE 'a%o'	Finds any values that start with "a" and end with "o"

SQL:

The screenshot shows a database interface with a query editor and a results grid. The query in the editor is:

```
SELECT CustLastName FROM Customer  
WHERE CustLastName LIKE "Sm%"
```

The results grid displays the following data:

CustLastName
Smith
Smithies
Smythe
Smart



Aggregate functions operate on the (sub)set of values in a column of a relation (table) and return a single value

- AVG()
 - Average value
- MIN()
 - Minimum value
- MAX()
 - Maximum value
- Plus others
 - <http://dev.mysql.com/doc/refman/8.0/en/group-by-functions.html>
- N.B. All of these except for COUNT() ignore null values and return null if all values are null. COUNT() counts the rows not the values and thus even if the value is NULL it is still counted.



DATA MANIPULATION LANGUAGE

- | | |
|---------|---------------------------------|
| COUNT() | - returns the number of records |
| AVG() | - average of the values |

Examples:

```
SELECT COUNT(CustomerID)  
FROM Customer;
```

= How many customers do we have
(cardinality)

```
SELECT AVG(OutstandingBalance)  
FROM Account;
```

= What is the average balance of
ALL ACCOUNTS

```
SELECT AVG(OutstandingBalance)  
FROM Account  
WHERE CustomerID= 1;
```

= What is the average balance of
Accounts of Customer 1

```
SELECT AVG(OutstandingBalance)  
FROM Account  
GROUP BY CustomerID;
```

= What is the average balance
PER CUSTOMER



- **Group by** groups all records together over a set of attributes
- Frequently used with aggregate functions
- **Example:**

*What is the average balance **PER CUSTOMER**?*

```
SELECT AVG(OutstandingBalance)
FROM Account
GROUP BY CustomerID;
```

- The only way to put a selection condition over a group by statement is by using **having** clause
- **Example:**

What is the exact average balance per customer for customers whose average balance is under 10000?

```
SELECT AVG(OutstandingBalance)
FROM Account
GROUP BY CustomerID
HAVING AVG(OutstandingBalance) < 10000
```



Column renaming

We can rename the column name of the output by using the AS clause

```
• SELECT CustType, Count(CustomerID)
  FROM Customer
  GROUP BY CustType;
```

CustType	Count(CustomerID)
Personal	3
Company	6

```
• SELECT CustType, Count(CustomerID) AS Count
  FROM Customer
  GROUP BY CustType;
```

CustType	Count
Personal	3
Company	6



- Orders records by particular column(s)

ORDER BY XXX ASC/DESC (ASC is default)

SQL

```
• SELECT CustLastName, CustType  
FROM Customer  
ORDER BY CustLastName;
```

CustLastName	CustType
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company
Unila	Company

RESULT

```
• SELECT CustLastName, CustType  
FROM Customer  
ORDER BY CustLastName DESC;
```

CustLastName	Cust Type
Unila	Company
Smythe	Company
Smithies	Company
Smith	Personal
Smart	Company
Samson	Personal
Lam	Personal
Jones	Company
Jones	Company



- LIMIT number
 - limits the output size
- OFFSET number
 - skips first 'number' records

```
SELECT CustLastName, CustType  
FROM Customer  
ORDER BY CustLastName  
LIMIT 5;
```

CustLastName	CustType
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company

```
SELECT CustLastName, CustType  
FROM Customer  
ORDER BY CustLastName  
LIMIT 5  
OFFSET 3;
```

CustLastName	CustType
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company

CustLastName	CustType
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company
Unila	Company





- **SELECT * FROM Rel1, Rel2;** - this is a **cross product**

```
SELECT * FROM Customer, Account;
```

										CustomerID
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID	
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1	
2	James	NULL	Jones	JJ Enterprises	Company	1	Peter Smith	245.25	1	
3	Akin	NULL	Smithies	Bay Wart	Company	1	Peter Smith	245.25	1	
1	Peter	NULL	Smith	NULL	Personal	2	JJ Ent.	552.39	2	
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	2	
3	Akin	NULL	Smithies	Bay Wart	Company	2	JJ Ent.	552.39	2	
1	Peter	NULL	Smith	NULL	Personal	3	JJ Ent. Mgr	10.25	2	
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	2	
3	Akin	NULL	Smithies	Bay Wart	Company	3	JJ Ent. Mgr	10.25	2	

Not very useful...

Typically we would like to find:

For every record in the Customer table list every record in the Account table



- **Inner/Equi join:**

- Joins the tables over keys

```
SELECT * FROM Customer INNER JOIN Account  
ON Customer.CustomerID = Account.CustomerID; CONDITION
```

The screenshot shows a database interface with a query window at the top and a results grid below. The query is:

```
SELECT * FROM Customer INNER JOIN Account  
ON Customer.CustomerID = Account.CustomerID;
```

The results grid has columns: CustomerID, CustFirstName, CustMiddleName, CustLastName, BusinessName, CustType, AccountID, AccountName, OutstandingBalance, CustomerID. The data is as follows:

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	2

- **Natural Join:**

- Joins the tables over keys. The condition does not have to be specified (natural join does it automatically), but key attributes have to have the same name.

```
SELECT * FROM Customer NATURAL JOIN Account;
```

The screenshot shows a database interface with a query window at the top and a results grid below. The query is:

```
SELECT * FROM Customer NATURAL JOIN Account;
```

The results grid has columns: CustomerID, CustFirstName, CustMiddleName, CustLastName, BusinessName, CustType, AccountID, AccountName, OutstandingBalance. The data is as follows:

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25



- Outer join:

- Joins the tables over keys
- Can be *left* or *right* (see difference below)
- Includes records that **don't match** the join from the other table

```
SELECT * FROM Customer LEFT OUTER JOIN Account  
ON Customer.CustomerID = Account.CustomerID;
```

Customer Data									
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENt.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENt. Mgr	10.25	2
3	Akin	NULL	Smithies	Bay Wart	Company	NULL	NULL	NULL	NULL

```
SELECT * FROM Customer RIGHT OUTER JOIN Account  
ON Customer.CustomerID = Account.CustomerID;
```

Customer Data									
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENt.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENt. Mgr	10.25	2



DATA MANAGEMENT

- You need to know how to write SQL
 - DDL
 - DML
 - SELECT



INFO90002 DB systems & Info. Modelling

- SQL Summary
 - Overview of concepts, more examples